


Gaetano Chiriaco 882638  
Gianmarco Russo 887277



# Pill Quality Control: Classification & GANs

# Problem Outline

## *Pill quality control dataset:*

1. 225x225px RGB .jpg images
2. 3 different classes: *normal*, *chip*, *dirt*
3. Small dataset: 330 images
4. Unbalanced classes

## *Chosen Tasks:*

1. Image Classification
2. Image Augmentation using classic transformations and GANs



Normal  
149 images




Chip  
43 images



Dirt  
138 images

# First Image classification

- Stratified splitting in Train (60%), Validation (20%) and Test (20%)
- Convolutional Neural Network architecture:



Callbacks: EarlyStopping (patience=5)

Loss: Categorical Cross-Entropy

Optimizer: Adam (LR = 0.0001)


$$\text{Loss} = - \sum_{i=1}^{\text{output size}} y_i \cdot \log \hat{y}_i$$

$$w_t = w_{t-1} - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$$

Other approaches used:

- **Data augmentation** using the aforementioned net (horizontal/vertical flip and random rotation)
- Fine tuning of an imageNet pre-trained network (**MobileNet**)

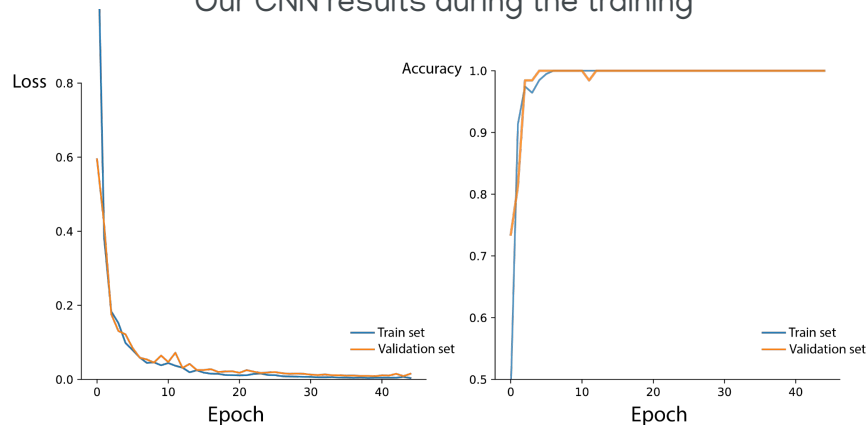
# Results on the original dataset



Our CNN does a perfect job on the test set

|  | F-Score | Accuracy | Precision | Recall | Loss   |
|--|---------|----------|-----------|--------|--------|
| Our CNN<br>390,083 trainable<br>parameters     | 1       | 1        | 1         | 1      | 0.0193 |
| MobileNet<br>1,865,283 trainable<br>parameters | 0.979   | 0.971    | 0.979     | 0.979  | 0.0359 |

Our CNN results during the training



When things go too well there's almost always something wrong...



# Fixing the biases - I

The models we trained have developed some biases: in particular the lighting and the background of the images is different among the classes



#6d8b8bff

#f3c4b0ff



#587175ff

#f7c6b8ff



#525753ff

#e7a390ff

Probably the model learns to distinguish between classes focusing on other aspects instead of the pill itself

# Fixing the biases - II


The training process has been repeated with **grayscale** and **masked** images.

The objective was to eliminate the biases caused by the **background** and the **lighting**.


**MobileNet** was not usable at this point since it only works with RGB images, so we only used our CNN, slightly tuning its architecture to work with grayscale images.



RGB




GrayScale



Without  
background



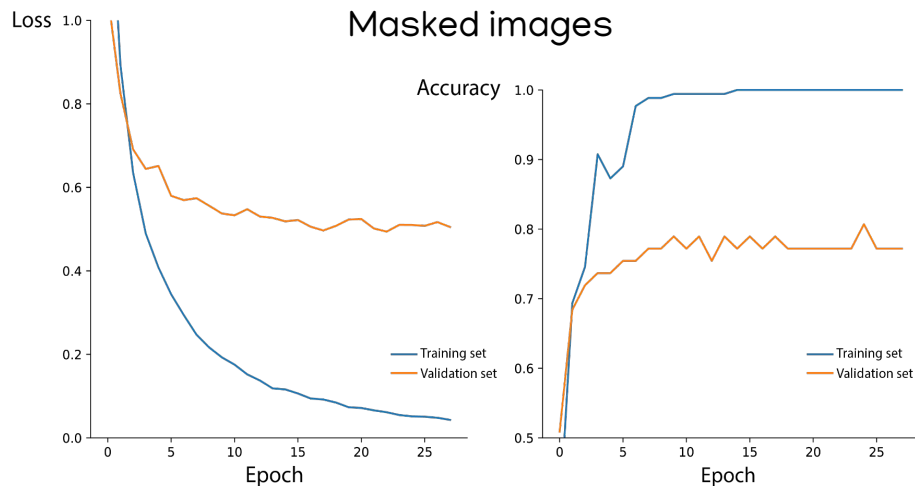

Not all the images were perfectly masked, so we removed 40 images that were badly cut



# Results on altered images

|                         | F-Score | Accuracy | Precision | Recall | Loss   | Class Accuracy                           |
|-------------------------|---------|----------|-----------|--------|--------|--|
| CNN on grayscale images | 0.989   | 0.9857   | 0.989     | 0.989  | 0.0568 | Normal: 1<br>Dirt: 0.965<br>Chip: 1      |
| CNN on masked images    | 0.769   | 0.787    | 0.835     | 0.718  | 0.5362 | Normal: 0.9<br>Dirt: 0.79<br>Chip: 0.375 |

As we can see, the grayscale images don't fix the biases, showing the same performance. However the masked one definitely showed some signs of change: the model can't exploit the background and the lighting and its performance drops quite significantly.

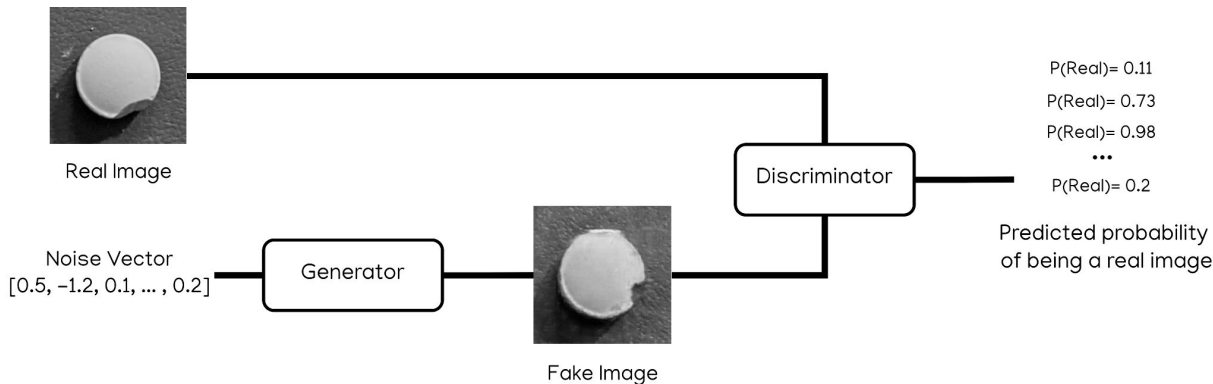


# Generative Adversarial Networks

**GANs** are generative models estimated via an **adversarial process**, in which we simultaneously train two models: a **generative model G** that captures the data distribution, and a **discriminative model D** that estimates the probability that a sample came from the training data rather than G.

Used for various tasks like:

- Image augmentation
- Super resolution
- Style transfer



Loss function: 
$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$



# GANs architecture for image augmentation

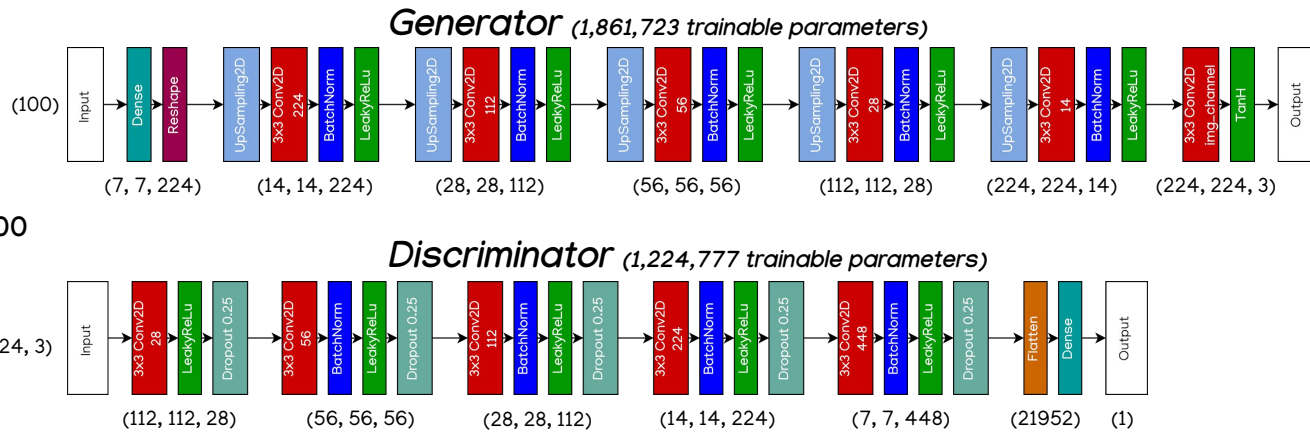
Different model for each pill type  
(normal, dirt, chip)

Different model for each image  
type (RGB, grayscale, masked)

500 epochs for the “normal” class, 500  
for the “dirt” class and 500 for the  
“chip” class


Noise size: 100

Batch size: 16




Problems and disadvantages encountered:


- Complicated structure
- Needs usually a high number of epochs to output good (fake) images
- Mode Collapse
- Hard to evaluate the effectiveness



# Data processing after GANs



# Comparing the performances



|   | F-Score | Accuracy | Precision | Recall | Loss   | Class Accuracy                           |
|---|---------|----------|-----------|--------|--------|--|
| CNN<br>on augmented<br>masked<br>images | 0.850   | 0.853    | 0.866     | 0.836  | 0.5326 | Normal: 0.76<br>Dirt: 0.92<br>Chip: 1    |
| CNN<br>on masked<br>images              | 0.769   | 0.787    | 0.835     | 0.718  | 0.5362 | Normal: 0.9<br>Dirt: 0.79<br>Chip: 0.375 |

In this kind of problems the accuracy on the “anomalous” classes is really what matters





# Possible developments and improvement

- Improve the masking process to avoid the removal of some of the images
- Explore different and more complicated types of GAN to obtain better results (like Conditional GANs)
- Solve the “Mode Collapse” issue
- Choose the best generation of images using more “formal” approaches (like Inception Score)

# Thanks for your attention!

**Pill Quality  
Control:  
Classification  
GANs**

Bibliography:

1.GANs: [\[1406.2661\]](https://arxiv.org/abs/1406.2661) Generative Adversarial Networks (arxiv.org)

Sitography:

1.Transfer learning with MobileNet: [https://www.tensorflow.org/tutorials/images/transfer\\_learning](https://www.tensorflow.org/tutorials/images/transfer_learning);

2.GAN | Generate Your Own Dataset using Generative Adversarial Networks ([analyticsvidhya.com](https://analyticsvidhya.com))